

# Tenuous Relations and Timeseries Analysis

Nathaniel Forde

Data Science @ Personio

# Agenda

1. Motivation: Background Problem
2. Structural Autoregressive Models
3. Structural Autoregressive Models in PyMC
  - Pause for Questions
4. Vector Autoregressive Models
5. Bayesian Vector Autoregressive Models in PyMC
6. Hierarchical BVARs in PyMC

# Motivation: Background Problem

# Disclaimers and Acknowledgements



## Note

No Personio data used in this presentation. The problem discussed motivated the research, but the models presented below will use fake or public data.

We will motivate the work with a discussion of App performance measurements and customer feedback metrics.

How do they influence one another and how to best distil the relationship between the two?



## Note

Acknowledgements: The work on the BVAR discussed here took inspiration and borrowed heavily from the PYMC labs blogpost [here](#). The example PRs benefited from the feedback of the PYMC devs.



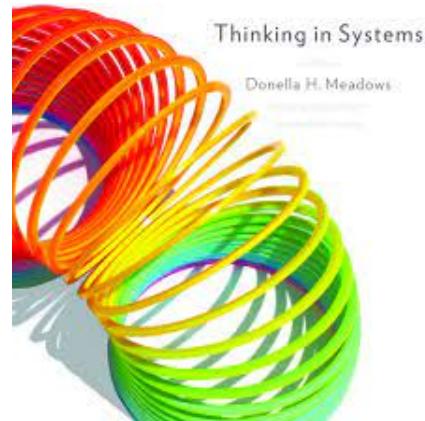
## Warning

Not an Economist!!!

# Systems and Feedback: App Performance

“Balancing feedback loops are equilibrating or goal seeking structures in systems and are both sources of stability and resistance to change”

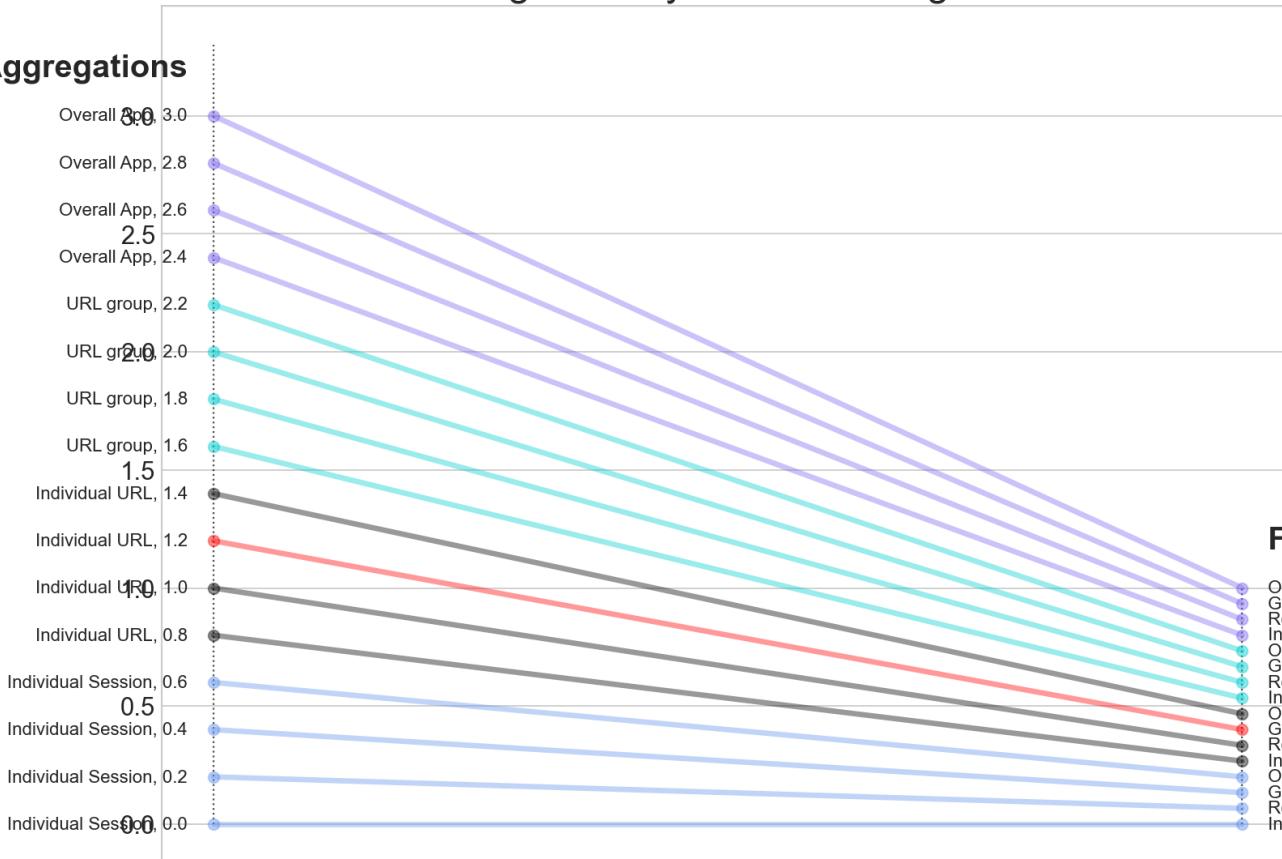
“Allowing performance standards to be influenced by past performance... sets up a reinforcing feedback of eroding goals that sets a system drifting toward low performance.”



# Disentangling Individual effects

Assessing the Ways feedback is generated

## App Performance Aggregations



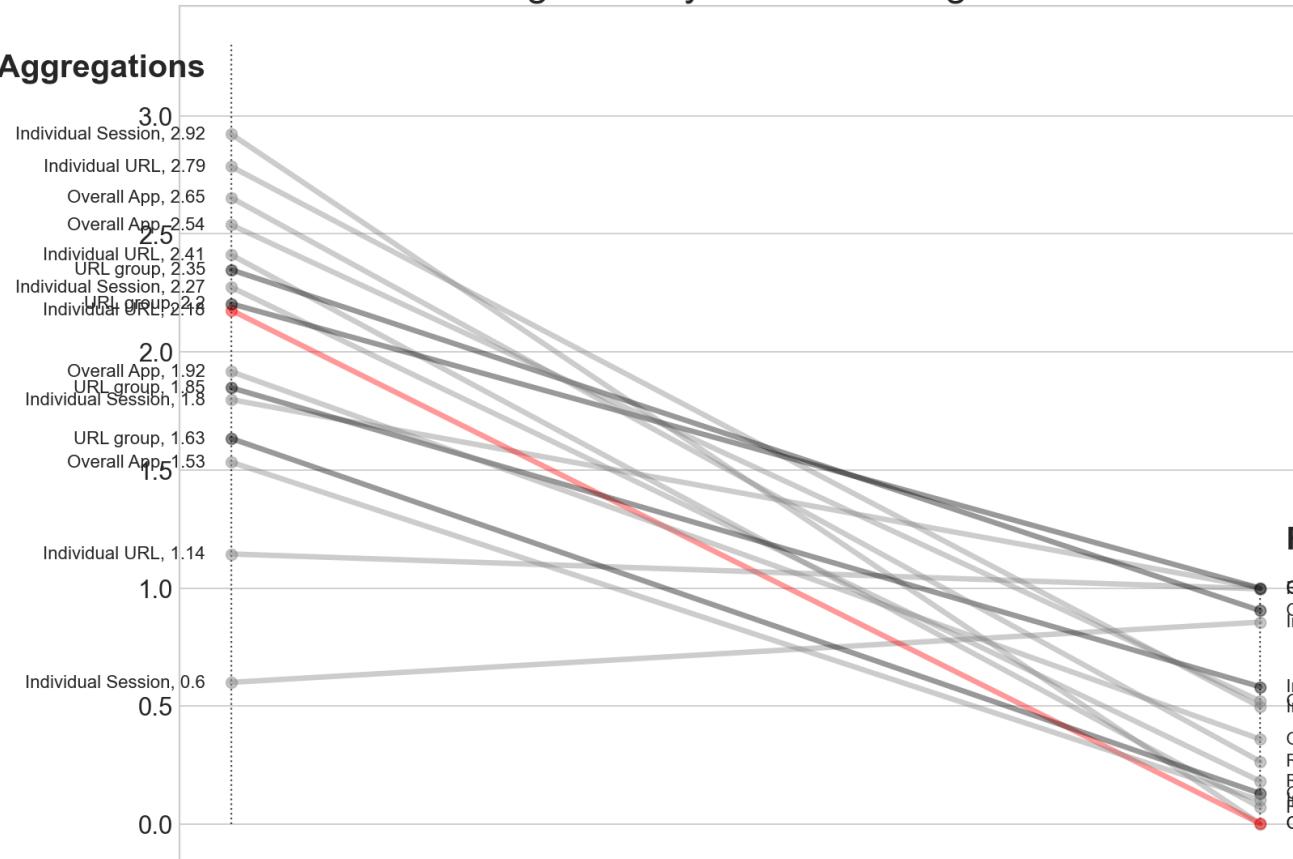
## Feedback Aggregations

- Overall Surveys, 1.0
- Grouped Surveys, 0.93
- Repeated Survey, 0.87
- Individual Survey, 0.8
- Overall Surveys, 0.73
- Grouped Surveys, 0.67
- Repeated Survey, 0.6
- Individual Survey, 0.53
- Overall Surveys, 0.47
- Grouped Surveys, 0.4
- Repeated Survey, 0.33
- Individual Survey, 0.27
- Overall Surveys, 0.2
- Grouped Surveys, 0.13
- Repeated Survey, 0.07
- Individual Survey, 0.0

# Identifying the Channel of Transmission

Assessing the Ways feedback is generated

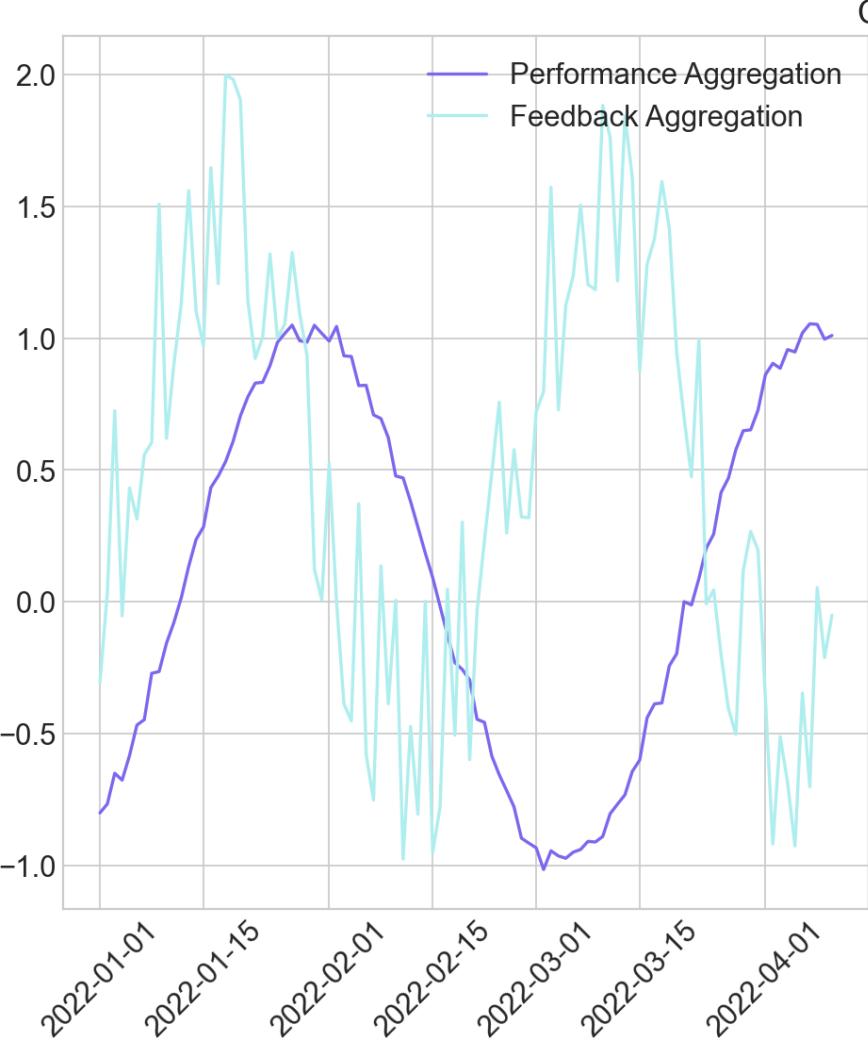
## App Performance Aggregations



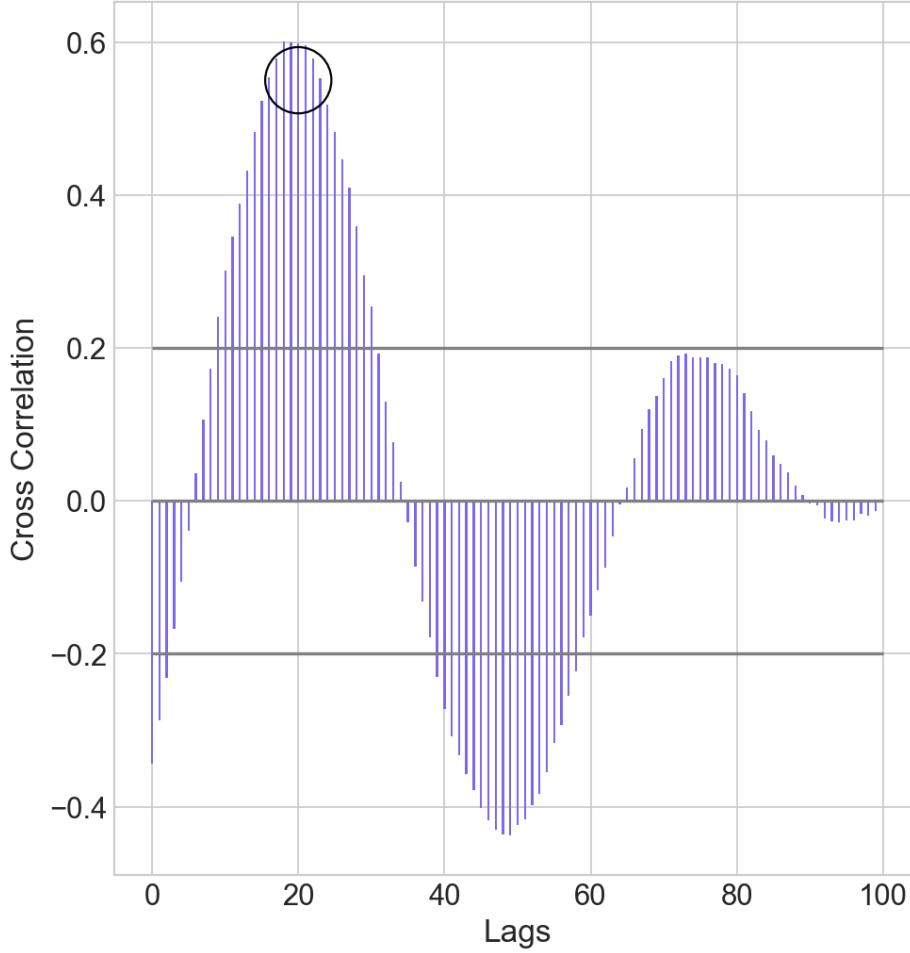
## Feedback Aggregations

Overall Surveys, 1.00
Overall Surveys, 0.9
Individual Survey, 0.86
Individual Survey, 0.58
Grouped Survey, 0.52
Overall Surveys, 0.36
Repeated Survey, 0.26
Repeated Survey, 0.18
Grouped Survey, 0.15
Repeated Survey, 0.07
Grouped Surveys, 0.0

# Lagging Effects in Time



Cross Correlation of Lagged terms between Performance and Feedback



# Structural Autoregressive Models

# Autoregressive Models

- Regression: The Linear Projection of  $Y$  onto  $X$  is the OLS solution.
- The Wold Decomposition: If a random variable  $Y_t$  is *covariance stationary* then  $Y_t$  has a linear representation:

$$Y_t = \mu + \sum_{j=0}^{\infty} b_j e_{t-j}$$

- Auto-regression: Extends this idea to Hilbert space of the infinite past history:

■

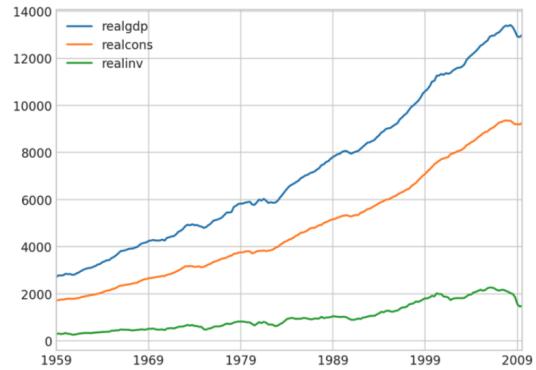
$$Y_t = \mu + \sum_{j=1}^{\infty} a_j Y_{t-j} + e_t$$

- Where the error terms  $e$  represent white-noise  $N(0, 1)$
- “...this justifies linear models as approximations” - Hansen *Econometrics*
- Most timeseries data is shorter than their infinite past.

# Moving Average Representation and Impulse Response

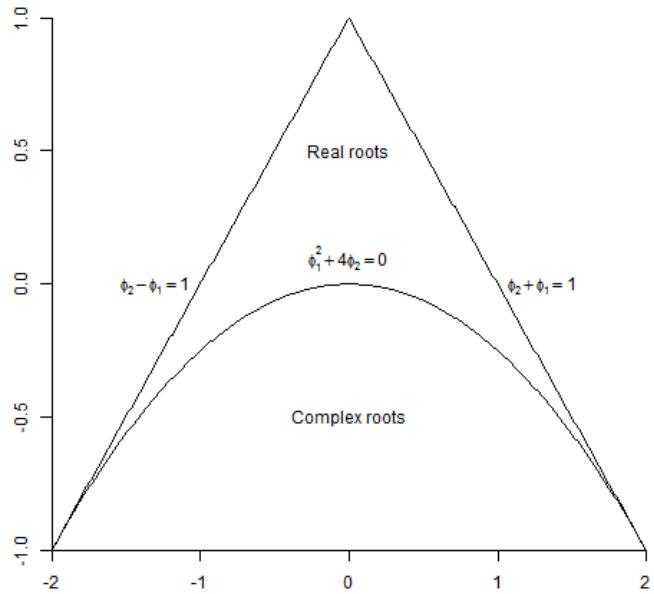
- The coefficients in:  $Y_t = \mu + \sum_{j=0}^{\infty} b_j e_{t-j}$  are known as the impulse response function **IRF** representing the change in  $Y_{t+j}$  due to a shock at time  $t$ .
  - $\frac{\partial}{\partial e_t} Y_{t+j} = b_j$
- They can be recovered from the AR representation by recursive reconstruction.
- IRFs are often reported when scaled by the standard deviation of shocks.

# Stationarity Requirements



Tricks to achieve stationarity

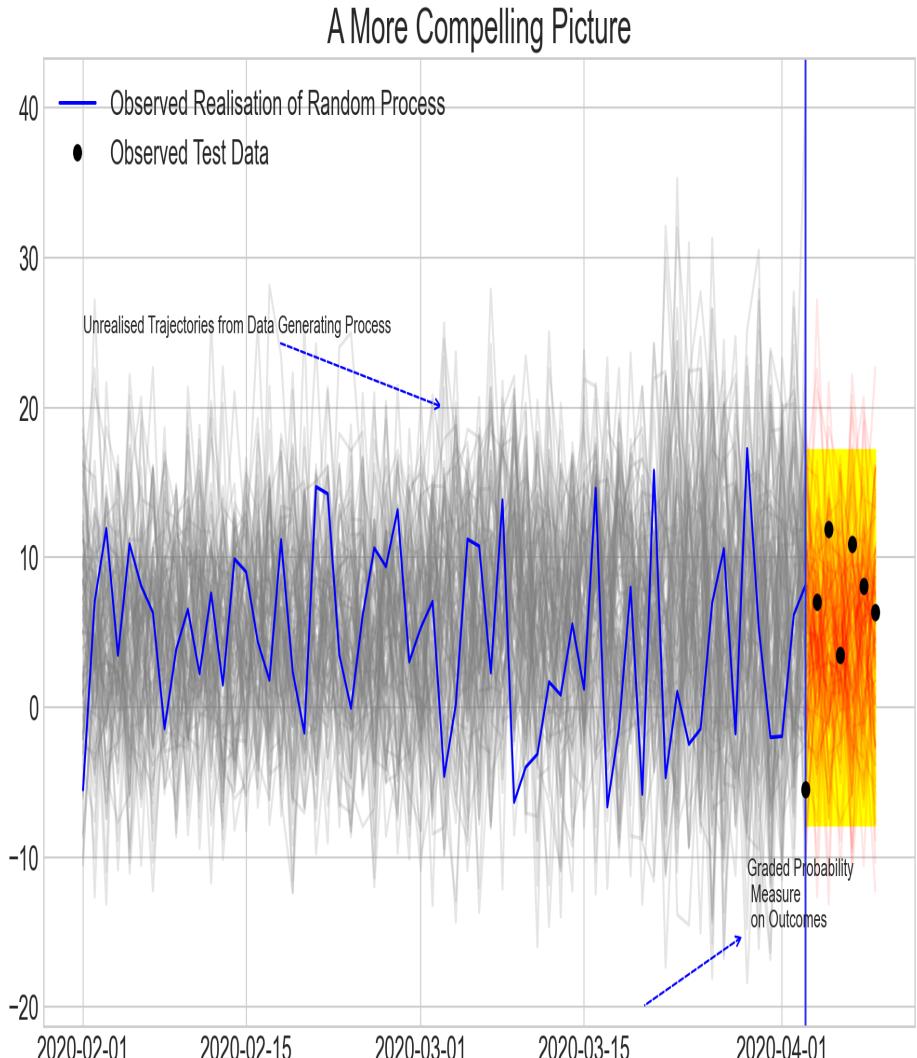
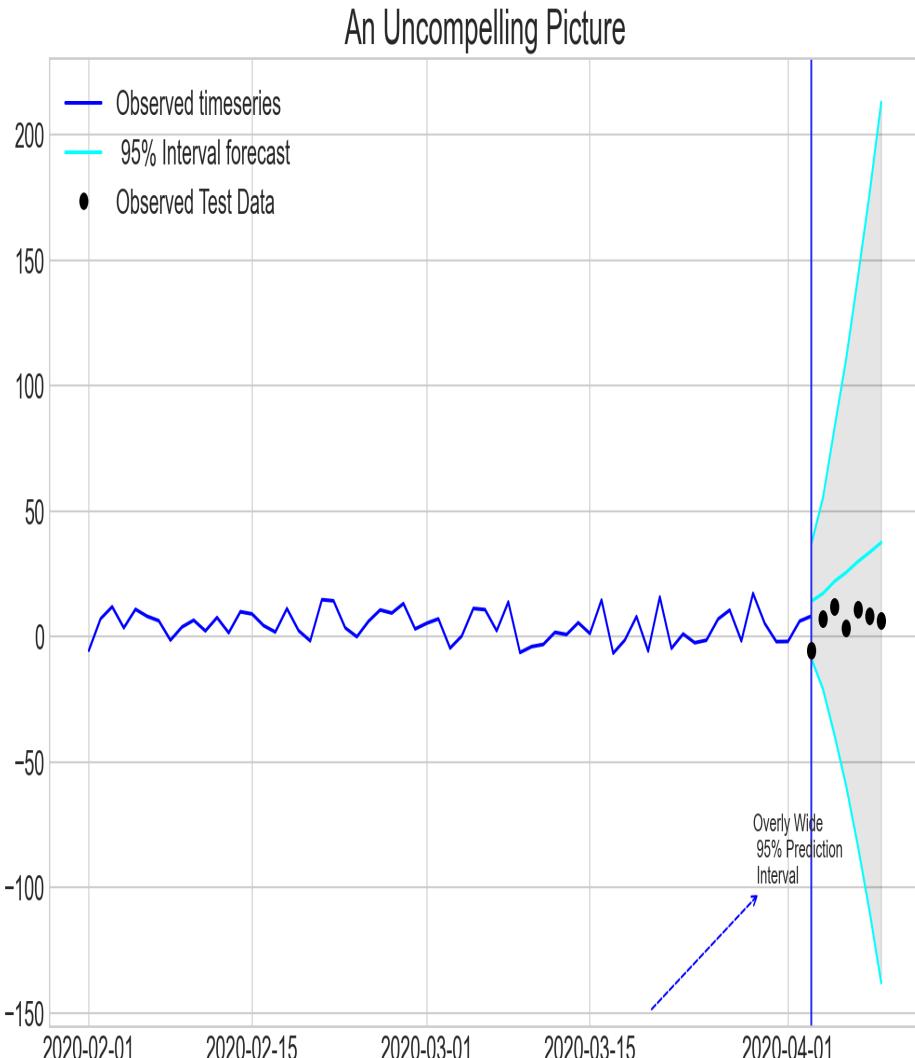
## Stationary conditions for AR(2) model



Conditions required for Stationarity

# Structure as Story

“The Child is the father of the Man” - Wordsworth



Personio

# Characterising Structure

$$Y \sim TruncNorm(\mu, \sigma, 0, \infty)$$

$$\mu = AR_\mu + f(seasonality) + f(trend) + \dots$$

$$\sigma \sim InverseGamma(3, .5)$$

$$trend = \alpha + \beta \cdot T_i$$

$$\alpha \sim Normal(7, 1)$$

$$\beta \sim Normal(0.009, .1)$$

$$seasonality = FourierFeatures' \beta_{fourier}$$

$$\beta_{fourier} \sim Normal(0, 0.5)$$

$$AR_\mu \sim \mu_{ar} + a_1 \cdot Y_{t-1} + a_2 \cdot Y_{t-2} \dots + \gamma$$

# Learning Structure

$$p(\theta | \text{Data}) \sim p(\theta) \cdot p(\text{Data} | \theta)$$

# Bayesian Structural Autoregressive Models

- Bayesian Structural Timeseries (BSTS) are often seen as an alternative to Auto-regressive models<sup>1</sup>
- But BSTS models can also incorporate latent auto-regressive components.
- Benefits of going Bayesian:
  - Transparent understanding of the sources of uncertainty within your model and a capacity to introduce outside information by way of priors or other covariates.
  - Expressive “lego-like” composability of the modelling structure matches the compositionality of thought.
  - Plausible counterfactual inference with posterior predictive distributions -

$$p(\hat{Data}|\mathbf{Data}) = \int_{\Theta} p(\hat{Data}|\theta, \mathbf{Data}) p(\theta|\mathbf{Data}) d\theta$$

<sup>1</sup> Nice post by Kim Larson [here](#)

# Structural Autoregressive Models in PyMC

# Bayesian Structural Timeseries

Priors

Additive Structural Components

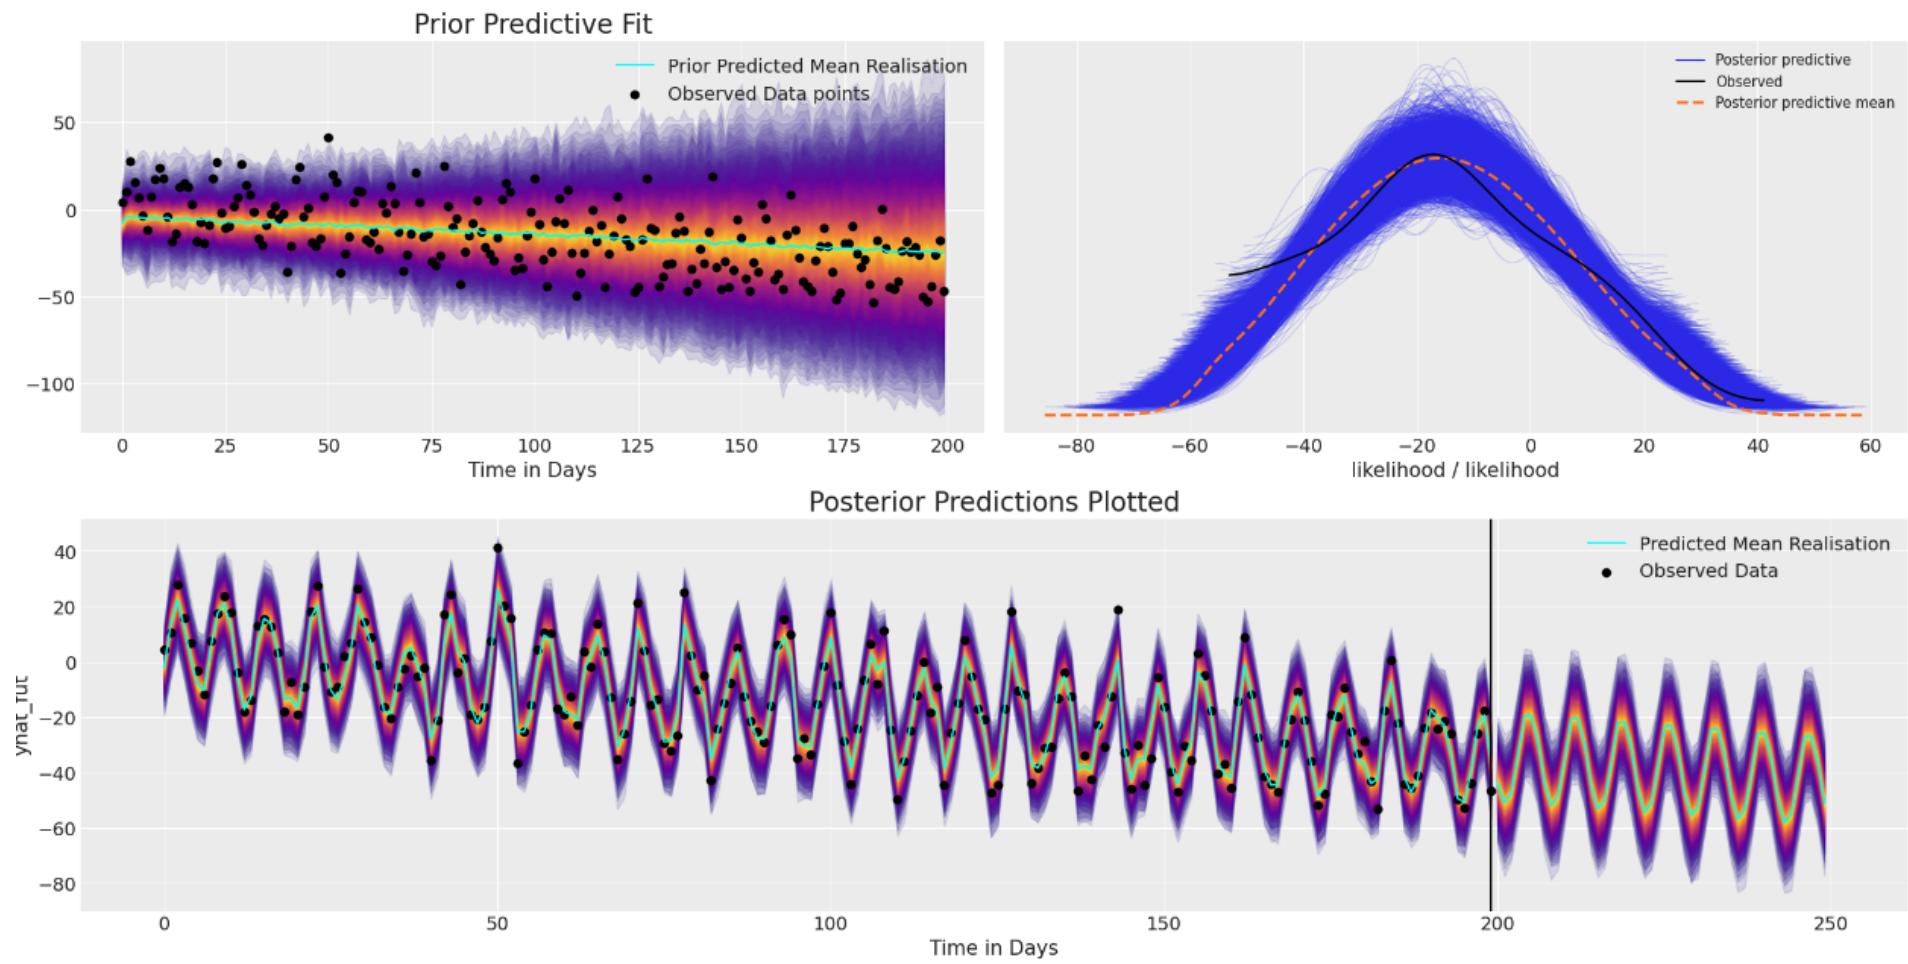
Likelihood

```

1
2 with AR:
3     ## Data containers to enable prediction
4     t = pm.MutableData("t", t_data, dims="obs_id")
5     y = pm.MutableData("y", ar_data, dims="obs_id")
6     # AR priors: The first coefficient will be the intercept term
7     coefs = pm.Normal("coefs", priors["coefs"]["mu"], priors["coefs"]["sigma"])
8     sigma = pm.HalfNormal("sigma", priors["sigma"])
9     ## Priors for the linear trend component
10    alpha = pm.Normal("alpha", priors["alpha"]["mu"], priors["alpha"]["sigma"])
11    beta = pm.Normal("beta", priors["beta"]["mu"], priors["beta"]["sigma"])
12    ## Priors for seasonality
13    beta_fourier = pm.Normal(
14        "beta_fourier",
15        mu=priors["beta_fourier"]["mu"],
16        sigma=priors["beta_fourier"]["sigma"],
17        dims="fourier_features",
18    )
19
20    # AR process: We need one init variable for each lag, hence size is variable too
21    init = pm.Normal.dist(
22        priors["init"]["mu"], priors["init"]["sigma"], size=priors["init"]["size"]
23    )
24    # Steps of the AR model minus the lags required given specification
25    ar1 = pm.AR(
26        "ar",
27        coefs,
28        sigma=sigma,
29        init_dist=init,

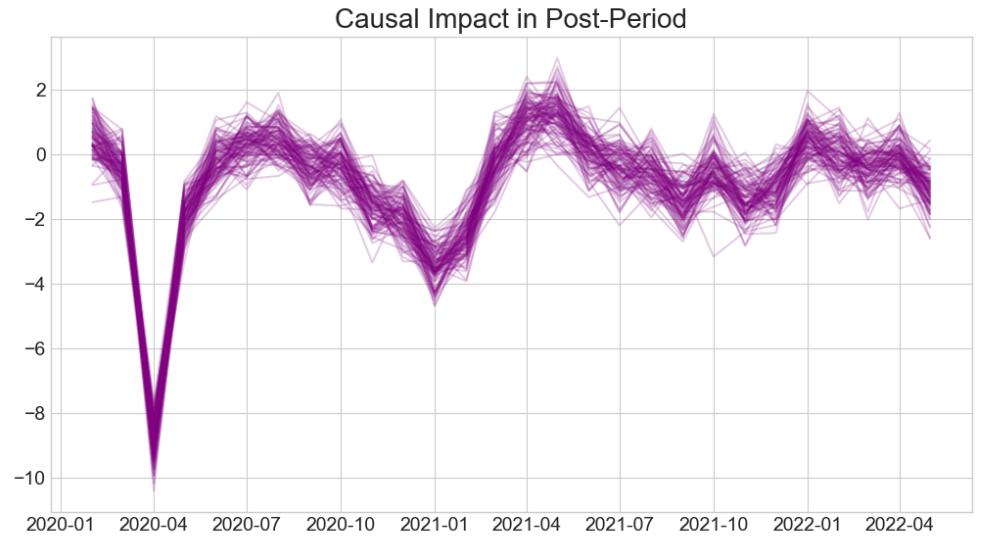
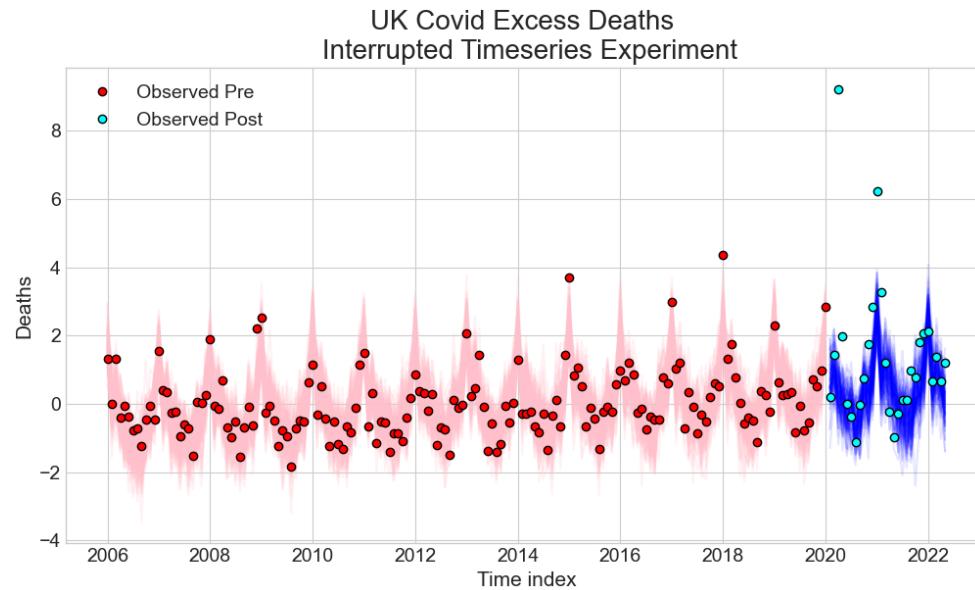
```

# Prior Predictive and Posterior Predictive Distributions



Autoregressive BSTS prior and posterior fits. For details see PYMC docs [here](#)

# Causal Inference with Interrupted Timeseries Designs



Causal Impact analysis for more detail see [CausalPy](#)

# Interim: Q & A

# Vector Autoregressive Models

# Multivariate Timeseries

- Multivariate Linear Projection Solution applies giving

- $\mathbf{y}_T = \nu + A_1 \mathbf{y}_{T-1} + A_2 \mathbf{y}_{T-2} \dots A_p \mathbf{y}_{T-p} + \mathbf{e}_t$

- Wold Decomposition:

$$\mathbf{y}_T = \mu + \Omega_1 \mathbf{e}_{T-1} + \Omega_2 \mathbf{e}_{T-2} \dots \Omega_p \mathbf{e}_{T-p}$$

- Multivariate Auto-regressive representation with MV white noise error terms

$$\begin{bmatrix} gdp \\ inv \\ con \end{bmatrix}_T = \nu + A_1 \begin{bmatrix} gdp \\ inv \\ con \end{bmatrix}_{T-1} + A_2 \begin{bmatrix} gdp \\ inv \\ con \end{bmatrix}_{T-2} \dots A_p \begin{bmatrix} gdp \\ inv \\ con \end{bmatrix}_{T-p} + \mathbf{e}_t$$

## VAR(2)

A VAR can have n lags and m equations such that the lagged terms of each series appear in all equations.

$$gdp_t = \beta_{gdp1} \cdot gdp_{t-1} + \beta_{gdp2} \cdot gdp_{t-2} + \beta_{cons1} \cdot cons_{t-1} + \beta_{cons2} \cdot cons_{t-2} + \epsilon_{gdp}$$

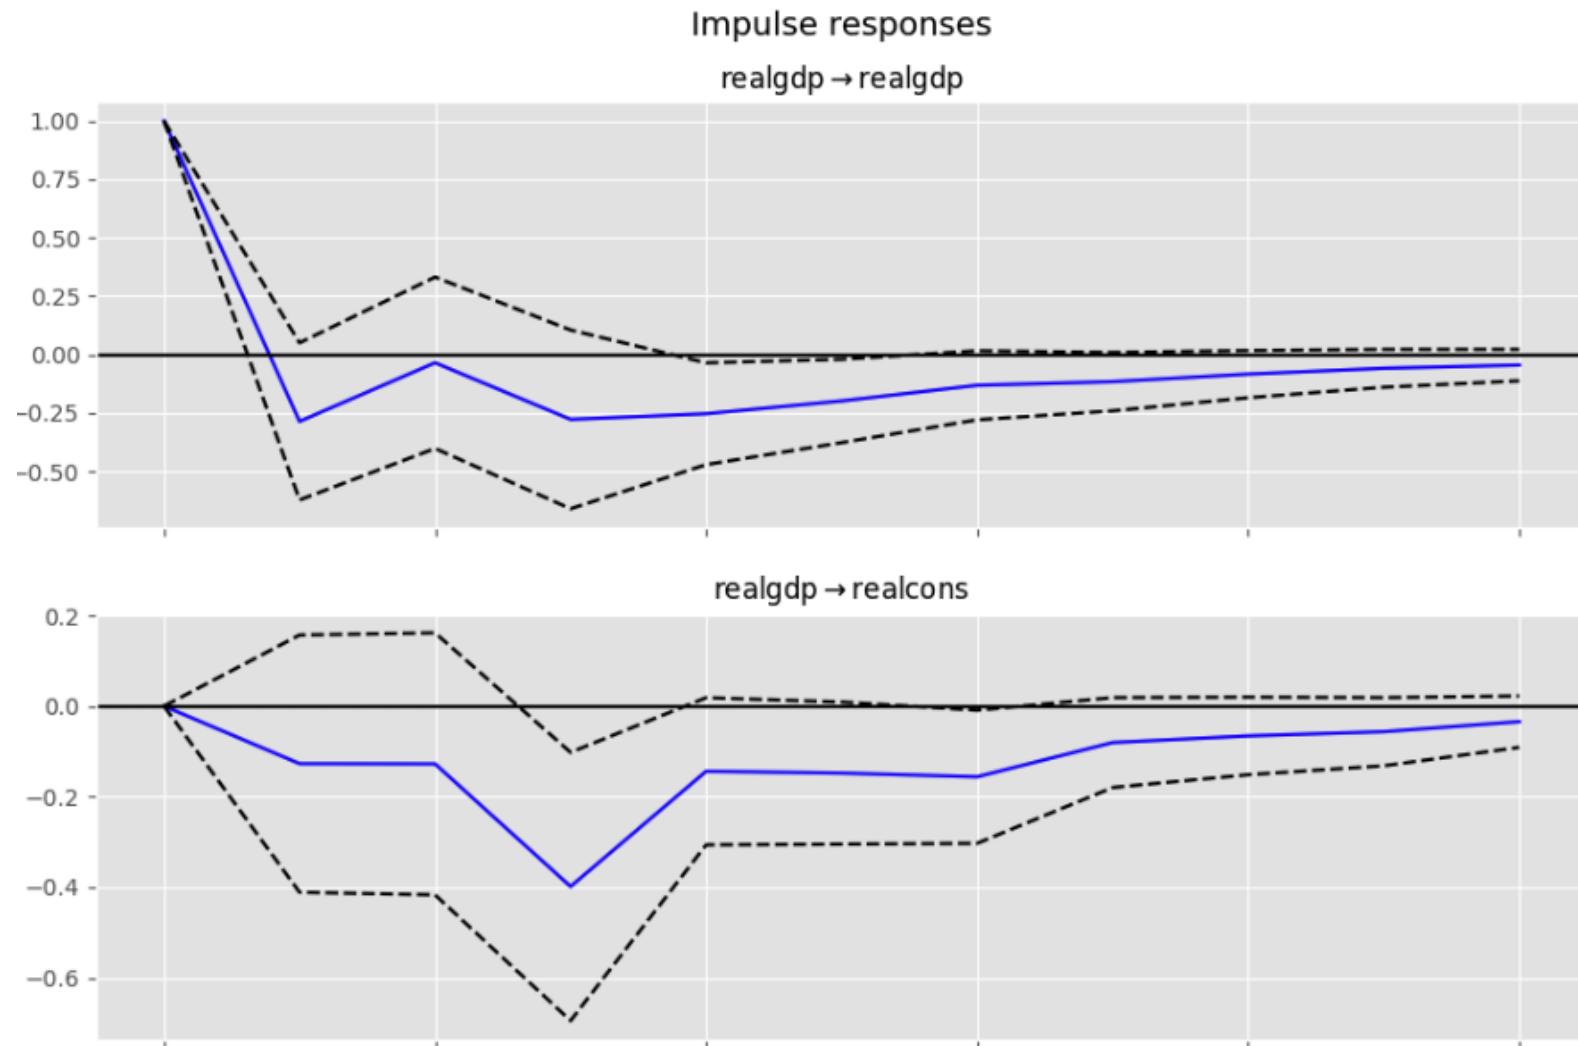
$$cons_t = \beta_{cons1} \cdot cons_{t-1} + \beta_{cons2} \cdot cons_{t-2} + \beta_{gdp1} \cdot gdp_{t-1} + \beta_{gdp2} \cdot gdp_{t-2} + \epsilon_{cons}$$

As such the **coefficients** on the cross-equation variables are of particular interest when investigating the influence of one series on another.

The MV IRF function can be created in an analogous way to the univariate timeseries function giving an interpretation of e.g the change in  $gdp_t$  due to a shock in  $cons_{t-2}$ .

This opens up the **possibility of testing** how influence between variables percolates over time.

# Impulse Response



Statsmodels implementation of IRF

# Vector Autoregressive Models in PyMC

# Arranging the Coefficient Matrices

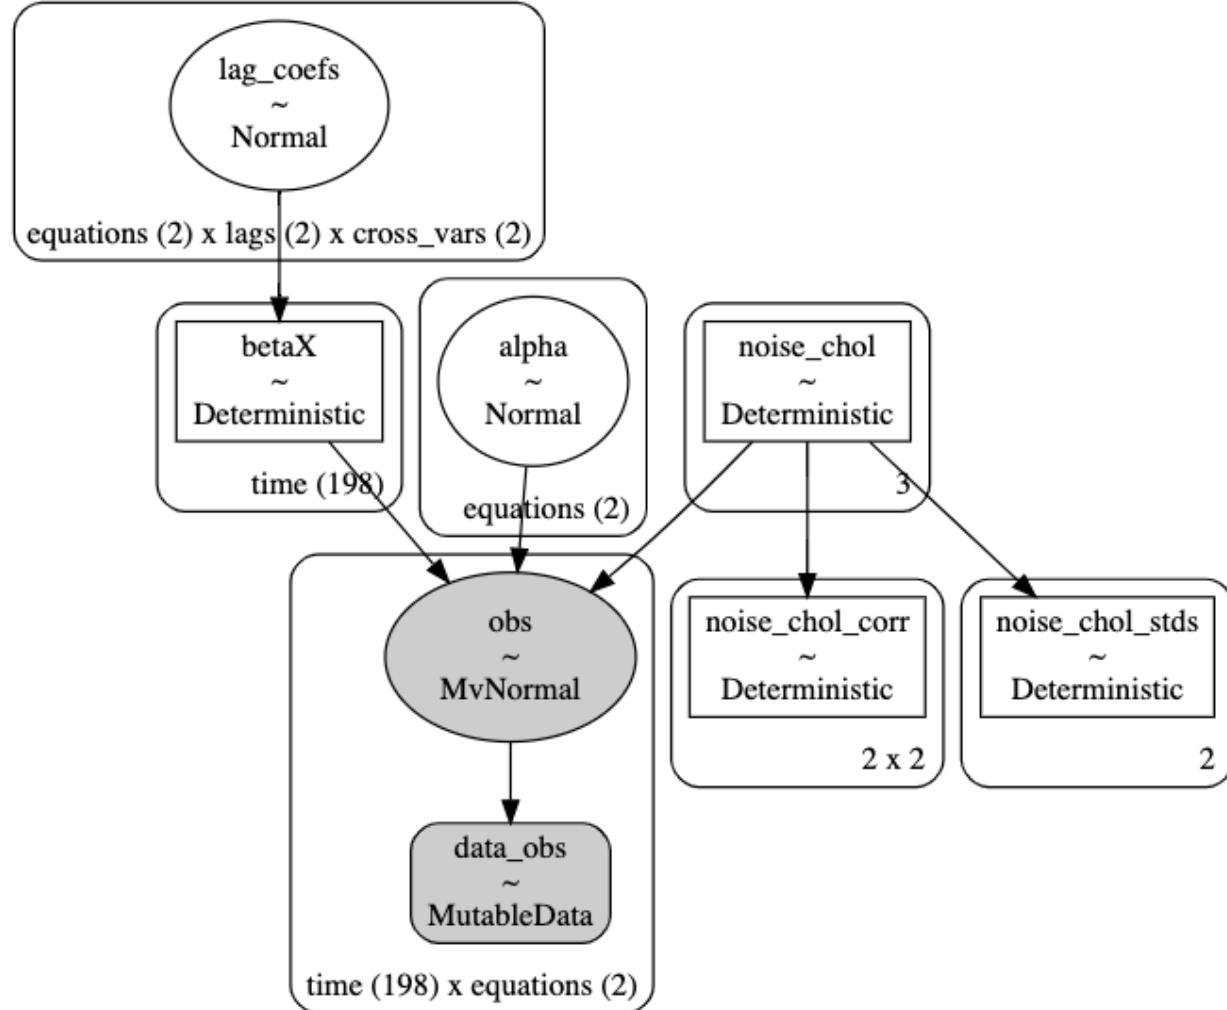
- There are allot of indices and variables in a VAR model.
- They can be concisely expressed in matrix notation, but it's not intuitive to keep track of all the indices.
- Good Bayesian workflow abstracts the component parts of the model definition.

```

1  ### Define a helper function that will construct our autoregressive step for the marginal contribution of each lagged
2  ### term in each of the respective time series equations
3  def calc_ar_step(lag_coefs, n_eqs, n_lags, df):
4      ars = []
5      for j in range(n_eqs):
6          ar = pm.math.sum(
7              [
8                  pm.math.sum(lag_coefs[j, i] * df.values[n_lags - (i + 1) : -(i + 1)], axis=-1)
9                  for i in range(n_lags)
10             ],
11             axis=0,
12         )
13     ars.append(ar)
14   betaX = pm.math.stack(ars, axis=-1)
15
16   return betaX

```

# The Broad Structure



VAR(2)

# Implementation in Code

```

1  with pm.Model(coords=coords) as model:
2      lag_coefs = pm.Normal(
3          "lag_coefs",
4          mu=priors["lag_coefs"]["mu"],
5          sigma=priors["lag_coefs"]["sigma"],
6          dims=["equations", "lags", "cross_vars"],
7      )
8      alpha = pm.Normal(
9          "alpha", mu=priors["alpha"]["mu"], sigma=priors["alpha"]["sigma"], dims=("equations",)
10     )
11     data_obs = pm.Data("data_obs", df.values[n_lags:], dims=["time", "equations"], mutable=True)
12
13     betaX = calc_ar_step(lag_coefs, n_eqs, n_lags, df)
14     betaX = pm.Deterministic("betaX", betaX, dims=["time",])
15     mean = alpha + betaX
16
17     if mv_norm:
18         n = df.shape[1]
19         ## Under the hood the LKJ prior will retain the correlation matrix too.
20         noise_chol, _, _ = pm.LKJCholeskyCov(
21             "noise_chol",
22             eta=priors["noise_chol"]["eta"],
23             n=n,
24             sd_dist=pm.HalfNormal.dist(sigma=priors["noise_chol"]["sigma"]),
25         )
26         obs = pm.MvNormal(
27             "obs", mu=mean, chol=noise_chol, observed=data_obs, dims=["time", "equations"]
28         )

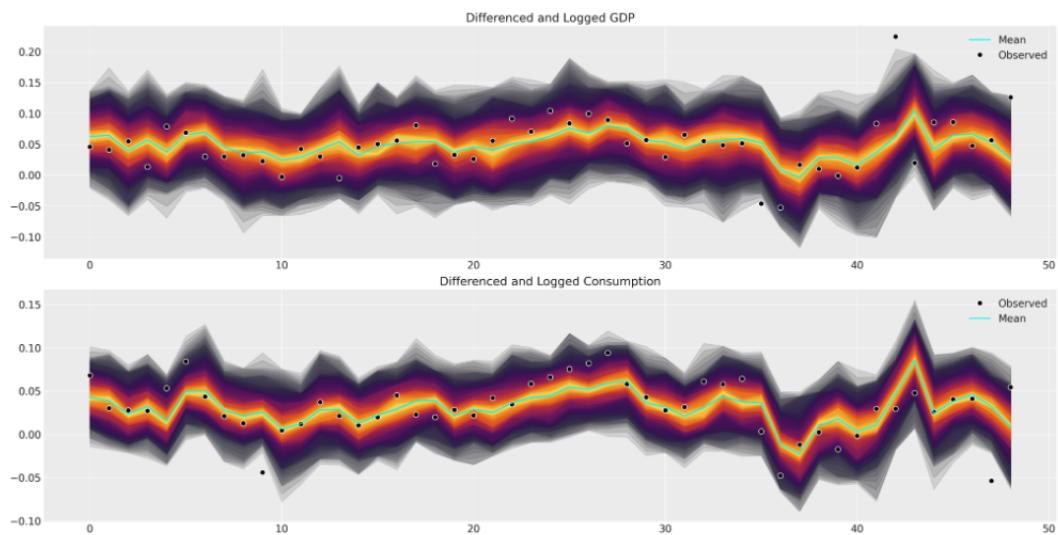
```

# Applying the Theory



Macroeconomic Data

# The Special Case of Ireland



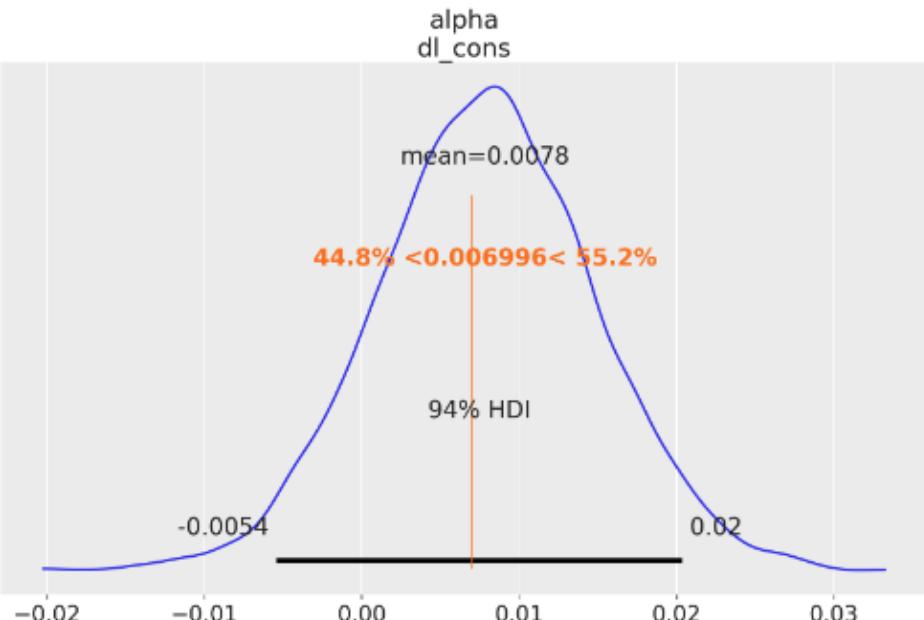
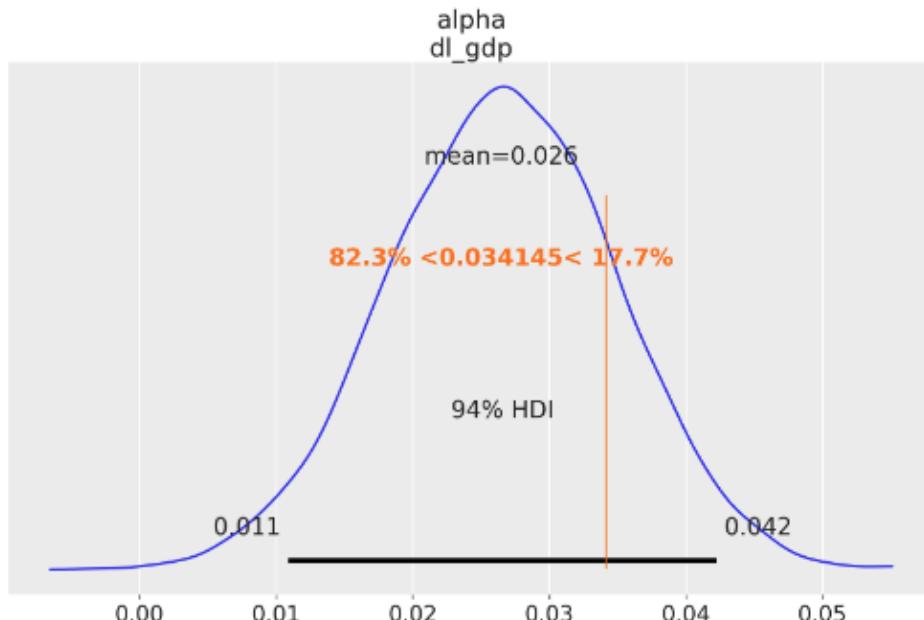
Ireland's Posterior Predictive Checks

Table 1: Posterior Mean Correlation

	GDP	CONS
GDP	1	0.43
CONS	0.43	1

# Comparison with Statsmodels MLE fit

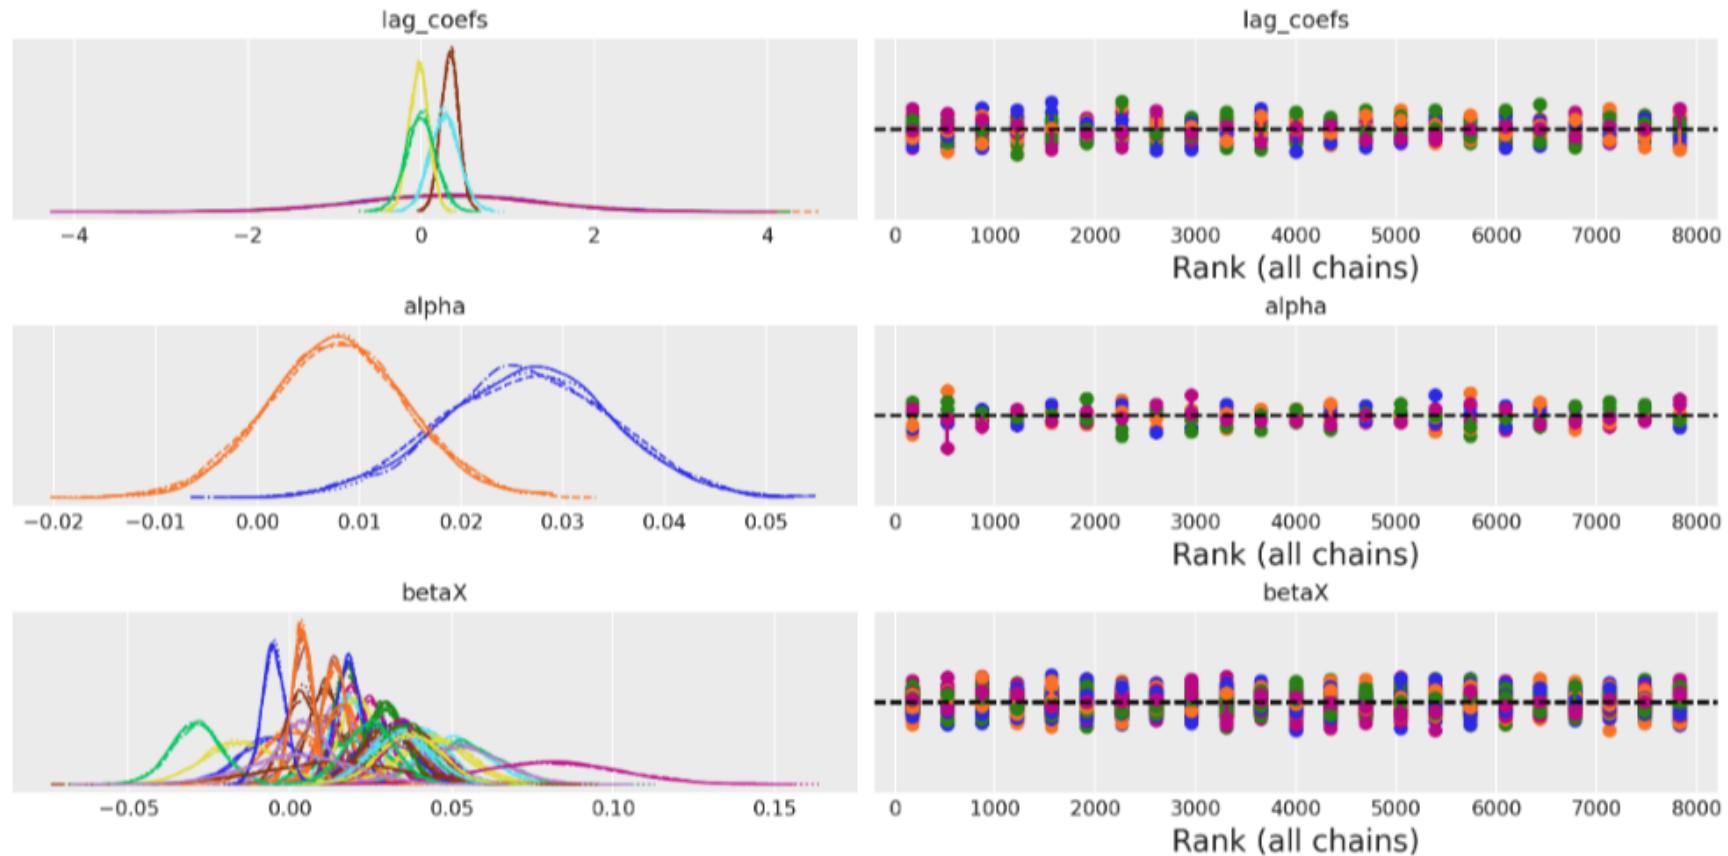
```
az.plot_posterior(idata_ireland, var_names=["alpha"], ref_val=[0.034145, 0.006996])
```



Compare

# Convergence Checks

```
az.plot_trace(idata_ireland, var_names=["lag_coefs", "alpha", "betaX"], kind="rank_
```

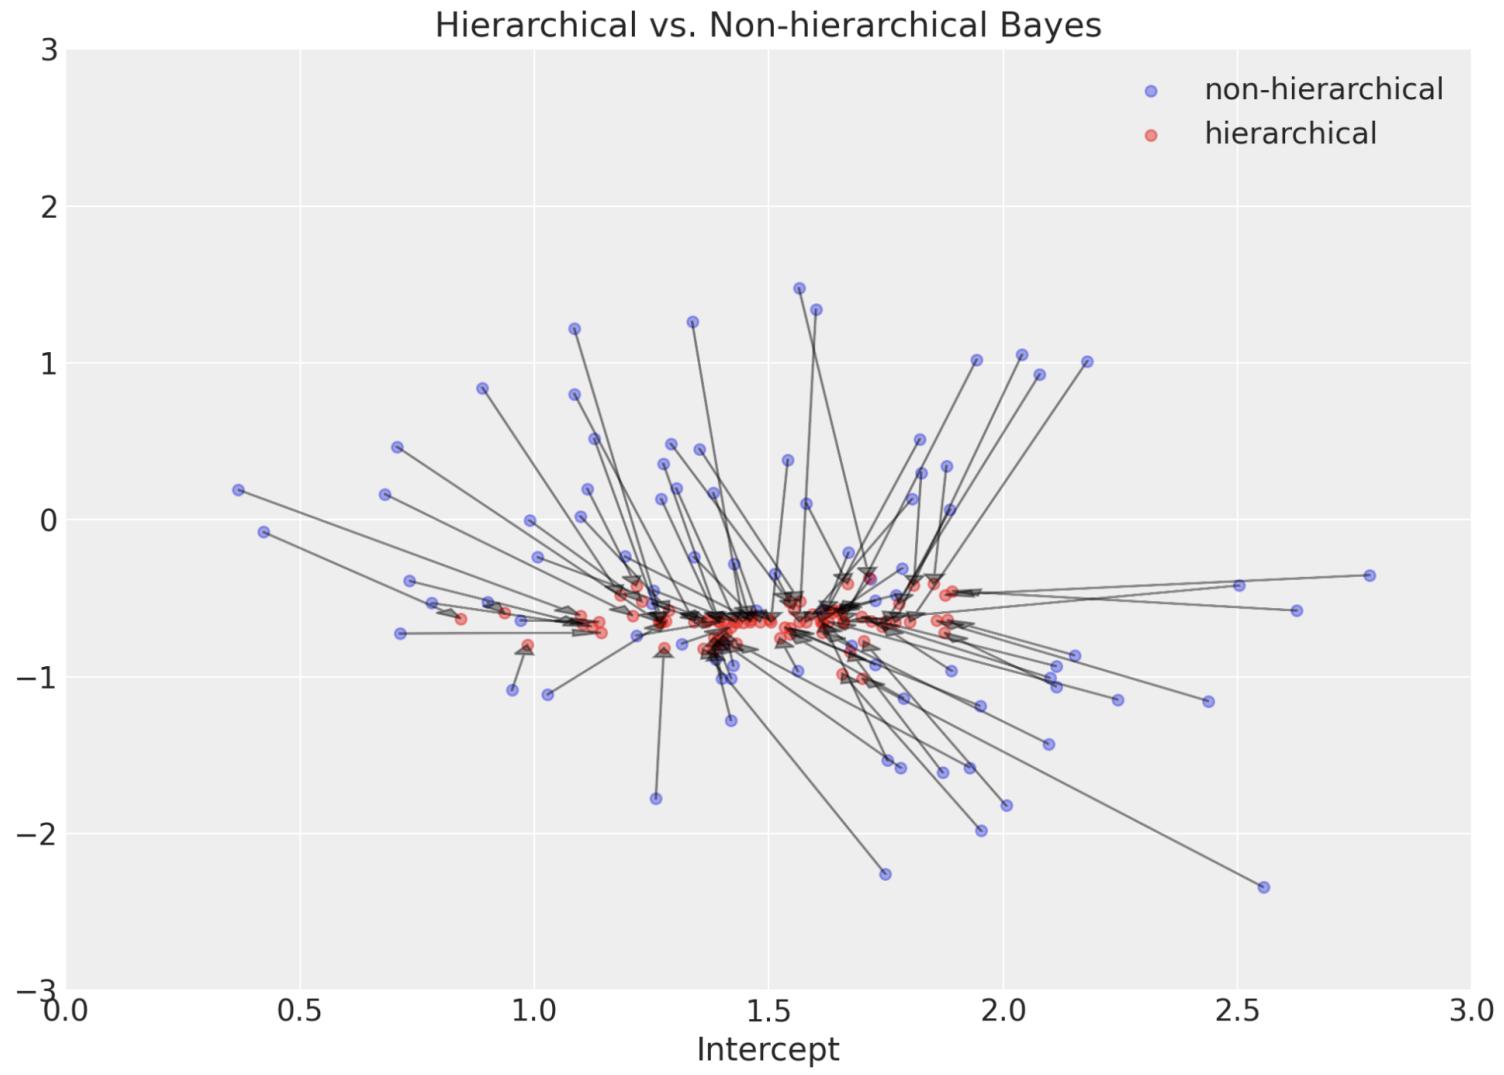


Sampling Chains

# Hierarchical Bayesian VARs

“Those who only know one country, know no country” – Seymour Martin Lipset

# Pooling, Shrinkage and Hierarchy



Shrinkage to Centre Effect

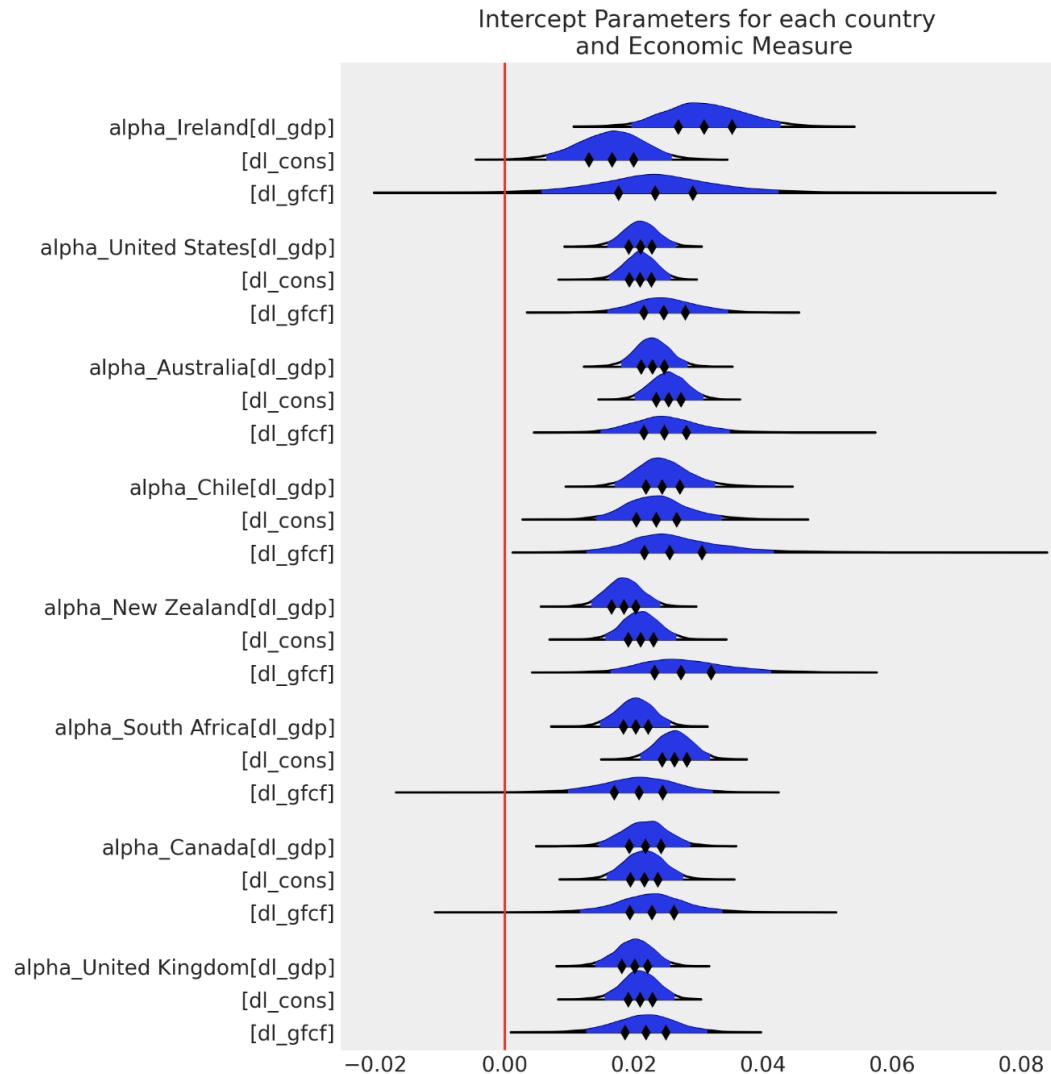
# How to model Hierarchical Structure

```

1 groups = df[group_field].unique()
2
3 with pm.Model(coords=coords) as model:
4     ## Hierarchical Priors
5     rho = pm.Beta("rho", alpha=2, beta=2)
6     alpha_hat_location = pm.Normal("alpha_hat_location", 0, 0.1)
7     alpha_hat_scale = pm.InverseGamma("alpha_hat_scale", 3, 0.5)
8     beta_hat_location = pm.Normal("beta_hat_location", 0, 0.1)
9     beta_hat_scale = pm.InverseGamma("beta_hat_scale", 3, 0.5)
10    omega_global, _, _ = pm.LKJCholeskyCov(
11        "omega_global", n=n_eqs, eta=1.0, sd_dist=pm.Exponential.dist(1)
12    )
13    ## Group Specific Parameter
14    for grp in groups:
15        df_grp = df[df[group_field] == grp][cols]
16        z_scale_beta = pm.InverseGamma(f"z_scale_beta_{grp}", 3, 0.5)
17        z_scale_alpha = pm.InverseGamma(f"z_scale_alpha_{grp}", 3, 0.5)
18        lag_coefs = pm.Normal(
19            f"lag_coefs_{grp}",
20            mu=beta_hat_location,
21            # Hierarchical offset
22            sigma=beta_hat_scale * z_scale_beta,
23            dims=["equations", "lags", "cross_vars"],
24        )
25        alpha = pm.Normal(
26            f"alpha_{grp}",
27            mu=alpha_hat_location,
28            # Hierarchical Offset
29            sigma=alpha_hat_scale * z_scale_alpha,

```

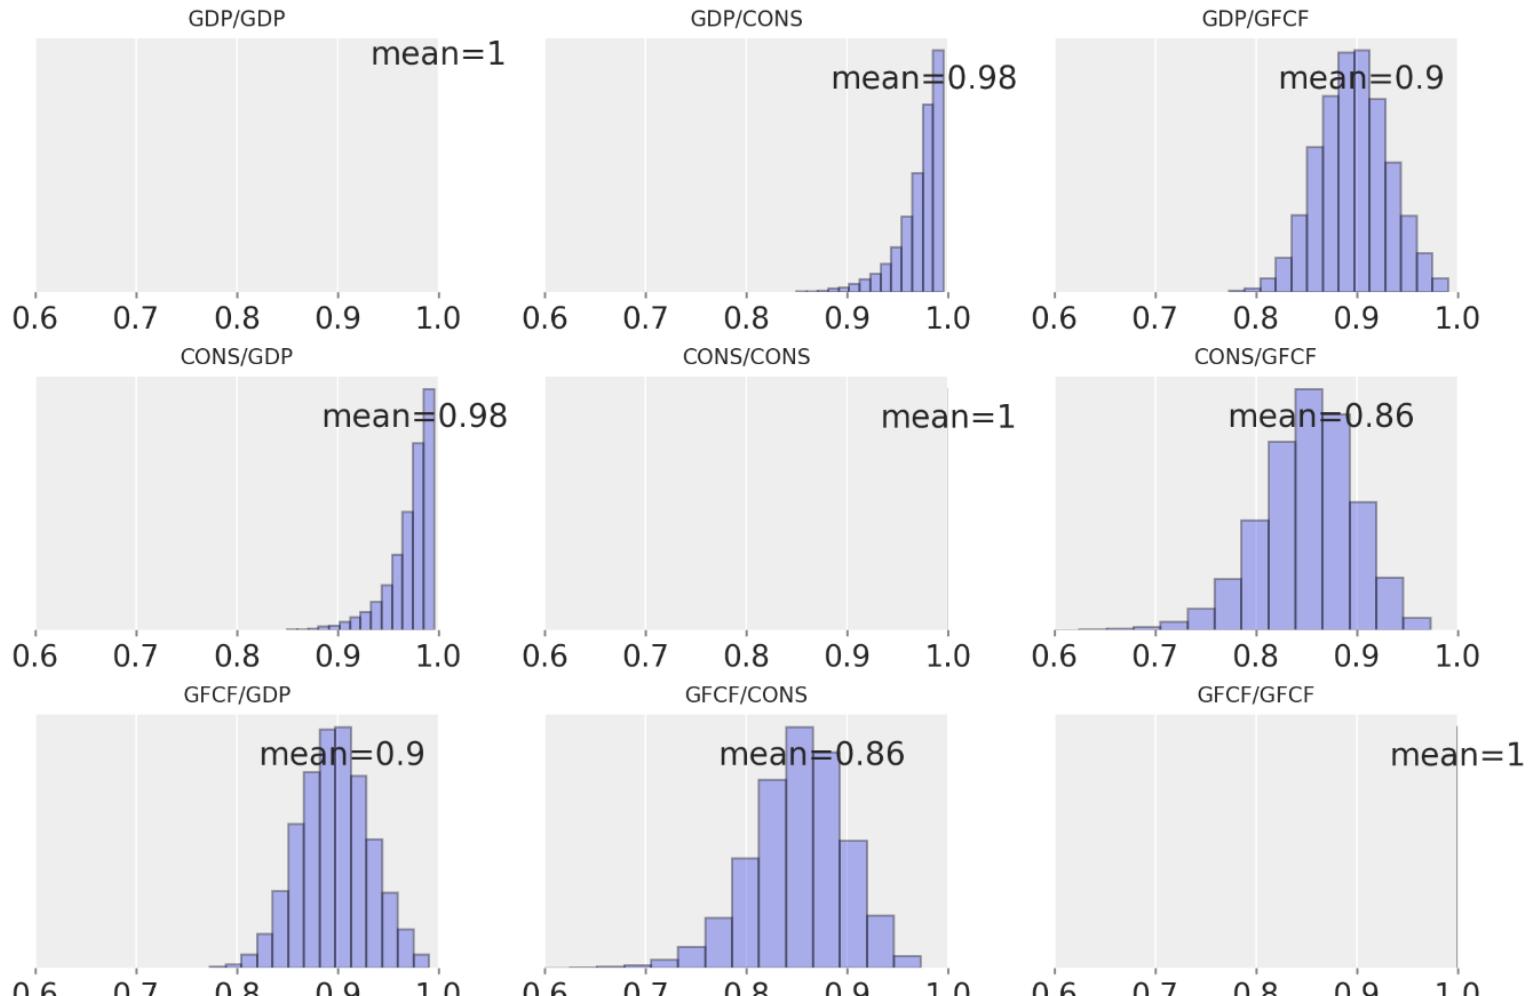
# Ireland is an Outlier



Country Specific Estimates

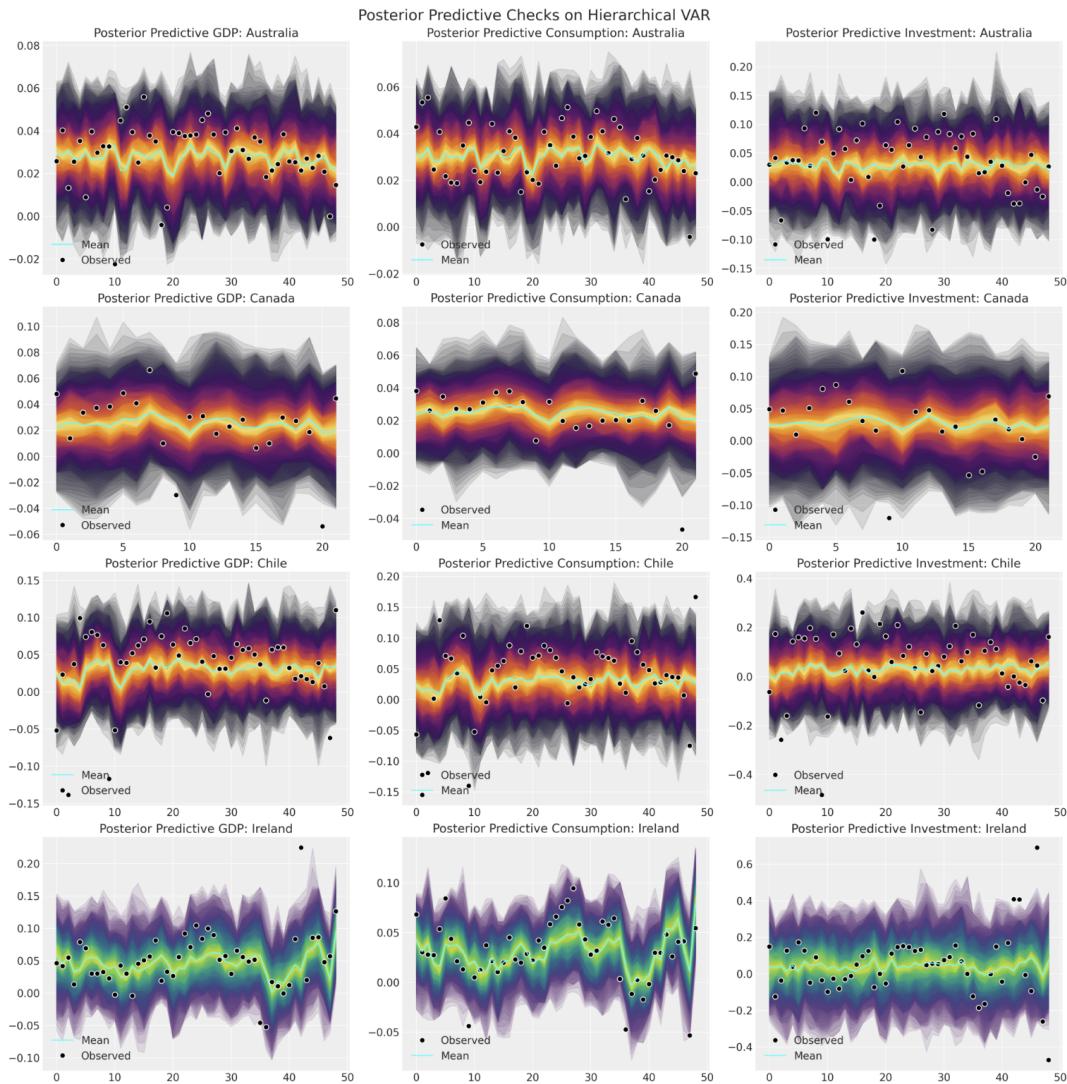
# The Global Correlation Structure

The Posterior Correlation Estimates



Global Correlations

# The Posterior Predictive Fits



Posterior Predictive Fits by country. See PYMC docs [here](#)

## Wrap Up

- Bayesian Timeseries models are flexible and transparent tools for forecasting and causal inference.
- VAR models can help interrogate questions of the relationships between timeseries data.
- Bayesian VAR models allow the specification of regularising priors and encode structural assumptions about direction of influence.
- Hierarchical Bayesian VAR models offer the chance to borrow information across groups and interrogate the questions about the relationships between timeseries within and across the groups.
- In the context with sparse or short timeseries the hierarchical model can augment our understanding of the lagged relationships of influence.

## For more Bayesian Content

—

# The Power of Bayes in Industry

## Your Business Model is Your Data Generating Process

 Feb 9, 2023 (Thurs) 21:00 UTC  
(1pm PT • 4pm ET • 10pm CET)



[pymcon.com/events](https://pymcon.com/events)

PYMCN

*Personio*

Personio